

```
1 !sudo apt-get install swig # swig 설치
2 !pip install box2d gymnasium[box2d]
```

 숨겨진 출력 표시

```
1 !git clone https://github.com/Soggeum/20242R0136C0SE47402.git
2 %cd 20242R0136C0SE47402/FinalProject/lunar-lander
```


 숨겨진 출력 표시

```
1 import gymnasium as gym
2 import pygame
3 import numpy as np
4 import matplotlib.pyplot as plt
5 import itertools
```

```
1 import agent_class as agent
```

✓ Initialize environment and agent


```
1 # We first create the environment on which we will later train the agent
2 env = gym.make('LunarLander-v3')
3
4 # We need to know the dimensionality of the state space, as well as how many
5 # actions are possible
6 N_actions = env.action_space.n
7 observation, info = env.reset()
8 N_state = len(observation)
9
10 print('dimension of state space =', N_state)
11 print('number of actions =', N_actions)
```

 dimension of state space = 8
number of actions = 4

```
1 # We create an instance of the agent class.
2 # At initialization, we need to provide
3 # - the dimensionality of the state space, as well as
4 # - the number of possible actions
5
6 parameters = {'N_state': N_state, 'N_actions': N_actions}
7
8 my_agent = agent.dqn(parameters=parameters)
9 # to train via the actor-critic algorithm, use this line:
10 # my_agent = agent.actor_critic(parameters=parameters)
```

✓ Train agent

```
1 # We train the agent on the LunarLander-v3 environment.
2 # Setting verbose=True allows us to follow the progress of the training
3
4 training_results = my_agent.train(environment=env,
5                                 verbose=True)
```



episode	return (this episode)	minimal return (last 20 episodes)	mean return (last 20 episodes)
100	-63.745	-333.606	-133.189
200	-137.825	-247.144	-97.487
300	-48.600	-399.012	-79.365
400	64.748	-264.094	-57.653
500	-11.198	-108.507	-5.500
600	-50.319	-182.694	21.444
700	92.862	-234.553	-12.551
800	230.056	-98.974	178.076
900	266.220	-7.204	208.497
1000	244.970	-212.061	163.451
1053	201.835	200.804	238.148

```
1 # the method my_agent.train() from the previous cell returns a dictionary
2 # with training stats, namely:
3 # - duration of each episode during training,
4 # - return of each episode during training
5 # - the total number of training epochs at the end of each episode
6 # - the total number of steps simulated at the end of each episode
```

```

7
8 training_results.keys()

dict_keys(['episode_durations', 'episode_returns', 'n_training_epochs', 'n_steps_simulated', 'training_completed'])

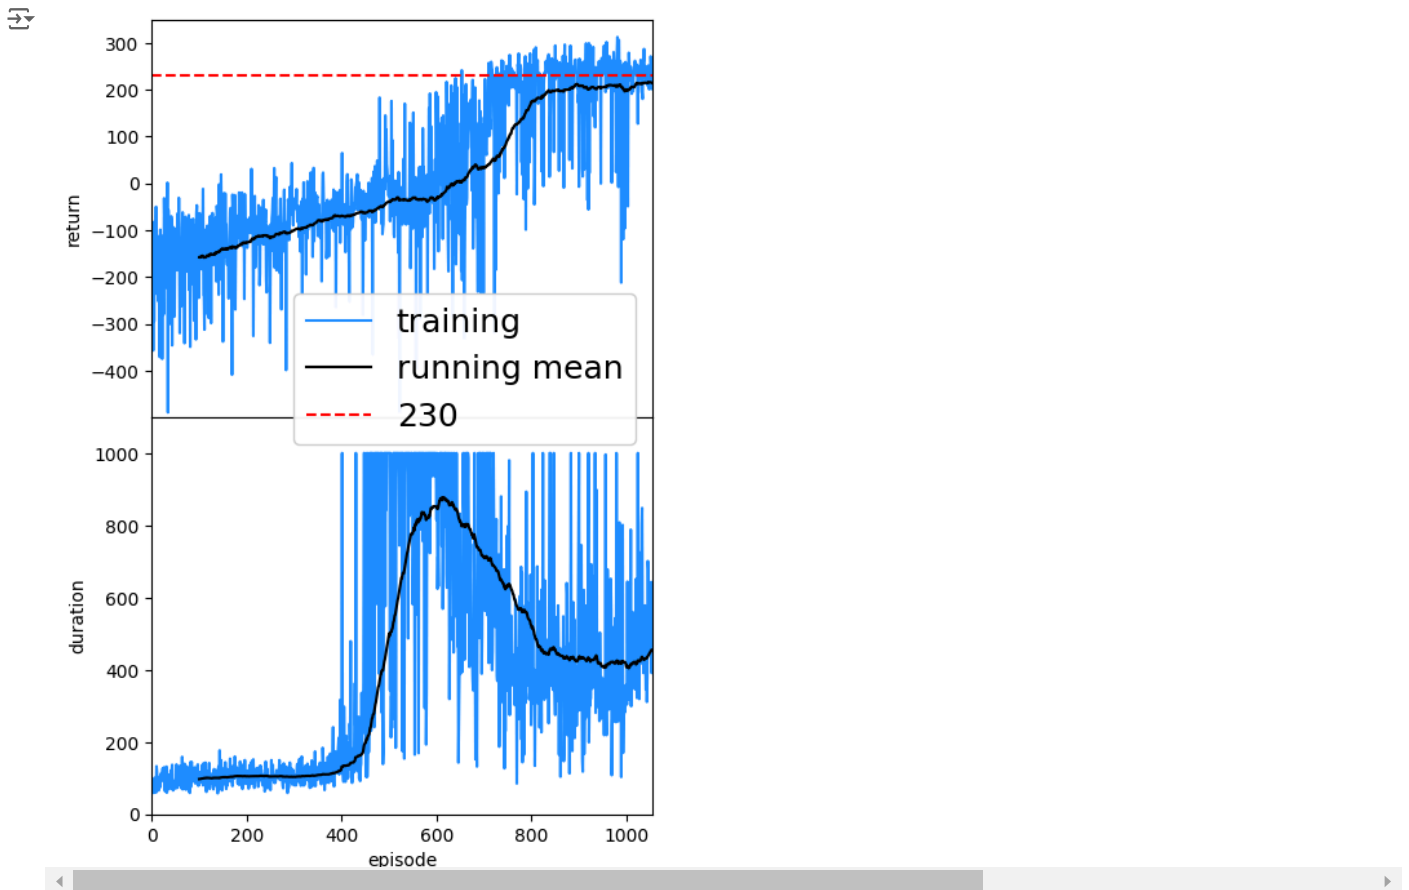
```

Plot training stats

```

1 # Plot both the return per episode and the duration per episode during
2 # training, together with their running average over 20 consecutive episodes
3
4 N = 100 # number of episodes for running average
5
6 def running_mean(x,N=100):
7     x_out = np.zeros(len(x)-N,dtype=float)
8     for i in range(len(x)-N):
9         x_out[i] = np.mean(x[i:i+N+1])
10    return x_out
11
12 def plot_returns_and_durations(training_results,filename=None):
13     fig,axes = plt.subplots(2,1,figsize=(5,8))
14     fig.subplots_adjust(hspace=0.0001)
15     #
16     # return as a function of episode
17     ax = axes[0]
18     x = training_results['episode_returns']
19     t = np.arange(len(x)) + 1
20     #
21     ax.plot(t,x,label='training',color='dodgerblue',)
22     # add running mean
23     x = running_mean(x=x,N=N)
24     t = np.arange(len(x)) + N
25     ax.plot(t,x,color='black',label='running mean')
26     #
27     ax.axhline(230,ls='--',
28               label='230',
29               color='red')
30     #
31     ax.set_ylim(-499,350)
32     ax.set_xticks([])
33     ax.set_xlim(0,len(t)+100)
34     ax.set_xlabel(r'episode')
35     ax.set_ylabel(r'return')
36     #
37     #
38     ax = axes[1]
39     x = training_results['episode_durations']
40     t = np.arange(len(x)) + 1
41     #
42     ax.plot(t,x,label='training',color='dodgerblue',)
43     # add running mean
44     x = running_mean(x=x,N=N)
45     t = np.arange(len(x)) + N
46     ax.plot(t,x,color='black',label='running mean')
47     #
48     ax.axhline(1200,ls='--', # draw line outside of plot scale,
49               label='230', # to get the red dotted line into the legend
50               color='red')
51     #
52     ax.set_ylim(0,1100)
53     ax.set_xlim(0,len(t)+100)
54     ax.set_xlabel(r'episode')
55     ax.set_ylabel(r'duration')
56     ax.legend(loc='upper right',bbox_to_anchor=(1.,1.35),
57              framealpha=0.95,
58              fontsize=18)
59     #
60     plt.show()
61     if filename != None:
62         fig.savefig(filename,bbox_inches='tight')
63     plt.close(fig)
64
65 plot_returns_and_durations(training_results=training_results)

```



✓ Create gameplay video using trained agent

First we create a "live" video that pops up and shows Lunar Lander gameplay performed by the agent

```
1 # There is the issue that the game window freezes when running gym games
2 # in jupyter notebooks, see https://github.com/openai/gym/issues/2433
3 # We here use the fix from that website, which is to use the following
4 # wrapper class:
5 class PyGameWrapper(gym.Wrapper):
6     def render(self, **kwargs):
7         retval = self.env.render( **kwargs)
8         for event in pygame.event.get():
9             pass
10        return retval
```

```
1 # Create a wrapped environment
2 env = PyGameWrapper(gym.make('LunarLander-v3',render_mode='human'))
3
4 N_episodes = 100
5
6 result_string = 'Run {0}: duration = {1}, total return = {2:7.3f}'
7
8 for j in range(N_episodes):
9     state, info = env.reset()
10
11     total_reward = 0
12     for i in itertools.count():
13         #env.render()
14
15         action = my_agent.act(state)
16         state, reward, terminated, truncated, info = env.step(action)
17         done = terminated or truncated
18         total_reward += reward
19
20     if done:
21         print(result_string.format(j+1,i+1,total_reward))
22         break
23
24 env.close()
```

```
Run 1: duration = 248, total return = 14.514
Run 2: duration = 377, total return = 248.985
Run 3: duration = 230, total return = 265.657
Run 4: duration = 369, total return = 191.265
Run 5: duration = 238, total return = 242.310
```

```

Run 6: duration = 269, total return = 239.245
Run 7: duration = 323, total return = 248.480
Run 8: duration = 381, total return = 218.104
Run 9: duration = 246, total return = 252.300
Run 10: duration = 470, total return = 221.463
Run 11: duration = 366, total return = 228.637
Run 12: duration = 264, total return = 225.442
Run 13: duration = 412, total return = 236.490
Run 14: duration = 397, total return = 226.231
Run 15: duration = 411, total return = 244.396
Run 16: duration = 277, total return = 256.117
Run 17: duration = 235, total return = 252.910
Run 18: duration = 465, total return = 257.295
Run 19: duration = 294, total return = 259.867
Run 20: duration = 242, total return = 242.040
Run 21: duration = 282, total return = 270.701
Run 22: duration = 318, total return = 260.515
Run 23: duration = 1000, total return = 133.132
Run 24: duration = 403, total return = 242.232
Run 25: duration = 500, total return = 248.794
Run 26: duration = 1000, total return = 146.760
Run 27: duration = 418, total return = 236.674
Run 28: duration = 297, total return = 224.078
Run 29: duration = 740, total return = 201.677
Run 30: duration = 277, total return = 261.403
Run 31: duration = 1000, total return = 144.660
Run 32: duration = 347, total return = 262.270
Run 33: duration = 1000, total return = 162.782
Run 34: duration = 267, total return = 37.562
Run 35: duration = 337, total return = 258.332
Run 36: duration = 1000, total return = 134.772
Run 37: duration = 344, total return = 254.130
Run 38: duration = 1000, total return = 105.749
Run 39: duration = 1000, total return = 143.841
Run 40: duration = 595, total return = 230.054
Run 41: duration = 322, total return = 241.073
Run 42: duration = 286, total return = 237.698
Run 43: duration = 1000, total return = 142.797
Run 44: duration = 408, total return = 252.662
Run 45: duration = 402, total return = 262.550
Run 46: duration = 319, total return = 230.427
Run 47: duration = 350, total return = 249.565
Run 48: duration = 316, total return = 222.992
Run 49: duration = 379, total return = 220.360
Run 50: duration = 493, total return = 278.496
Run 51: duration = 346, total return = 248.227
Run 52: duration = 485, total return = 248.675
Run 53: duration = 332, total return = 265.492
Run 54: duration = 582, total return = 234.948
Run 55: duration = 473, total return = 230.542
Run 56: duration = 1000, total return = 120.883
Run 57: duration = 389, total return = 252.505
Run 58: duration = 246, total return = 244.084

```

We also create a video file containing 20 games played by the agent

```

1 import gymnasium as gym
2 from gymnasium.wrappers import RecordEpisodeStatistics, RecordVideo
3
4 num_eval_episodes = 100
5
6 env = gym.make("LunarLander-v3", render_mode="rgb_array") # replace with your environment
7 env = RecordVideo(env, video_folder="my_video", name_prefix="eval",
8                   episode_trigger=lambda x: True)
9
10 env = RecordEpisodeStatistics(env, buffer_length=num_eval_episodes)
11
12 for episode_num in range(num_eval_episodes):
13     obs, info = env.reset()
14
15     episode_over = False
16     while not episode_over:
17         action = my_agent.act(state)
18         #action = env.action_space.sample() # replace with actual agent
19         obs, reward, terminated, truncated, info = env.step(action)
20
21         episode_over = terminated or truncated
22 env.close()
23
24 print(f'Episode time taken: {env.time_queue}')
25 print(f'Episode total rewards: {env.return_queue}')
26 print(f'Episode lengths: {env.length_queue}')

```

```

🔍 Episode time taken: deque([0.311925, 0.28265, 0.292123, 0.163196, 0.311078, 0.443475, 0.424414, 0.180685, 0.256797, 0.250661, 0.195404, 0.184452, 0.17
Episode total rewards: deque([-833.6724371955099, -737.155293017751, -769.1016327150238, -358.9979936134969, -499.80505331710754, -830.4024996506721,
Episode lengths: deque([87, 79, 85, 53, 63, 87, 74, 51, 77, 69, 57, 51, 51, 71, 87, 80, 64, 50, 53, 51, 56, 58, 62, 52, 64, 78, 51, 71, 80, 62, 83, 75

```

