

```
1 !sudo apt-get install swig # swig 설치
2 !pip install box2d gymnasium[box2d]
```

 숨겨진 출력 표시

```
1 !git clone https://github.com/Soggeum/20242R0136C0SE47402.git
2 %cd 20242R0136C0SE47402/FinalProject/lunar-lander
```


 숨겨진 출력 표시

```
1 import gymnasium as gym
2 import pygame
3 import numpy as np
4 import matplotlib.pyplot as plt
5 import itertools
```

```
1 import agent_class as agent
```

▽ Initialize environment and agent


```
1 # We first create the environment on which we will later train the agent
2 env = gym.make('LunarLander-v3')
3
4 # We need to know the dimensionality of the state space, as well as how many
5 # actions are possible
6 N_actions = env.action_space.n
7 observation, info = env.reset()
8 N_state = len(observation)
9
10 print('dimension of state space =', N_state)
11 print('number of actions =', N_actions)
```

 dimension of state space = 8
number of actions = 4

```
1 # We create an instance of the agent class.
2 # At initialization, we need to provide
3 # - the dimensionality of the state space, as well as
4 # - the number of possible actions
5
6 parameters = {'N_state': N_state, 'N_actions': N_actions}
7
8 my_agent = agent.dqn(parameters=parameters)
9 # to train via the actor-critic algorithm, use this line:
10 my_agent = agent.actor_critic(parameters=parameters)
```

▽ Train agent

```
1 # We train the agent on the LunarLander-v3 environment.
2 # Setting verbose=True allows us to follow the progress of the training
3
4 training_results = my_agent.train(environment=env,
5                                 verbose=True)
```



episode	return (this episode)	minimal return (last 20 episodes)	mean return (last 20 episodes)
100	-6.594	-371.114	-167.828
200	-88.411	-251.395	-136.390
300	-23.532	-213.534	-77.778
400	-216.276	-216.276	-54.640
500	-46.598	-168.718	-62.316
600	-4.206	-319.626	-36.105
700	-49.527	-128.284	73.911
800	-67.011	-148.418	104.010
900	271.473	-233.651	114.051
1000	262.370	-225.384	130.675
1100	302.630	-266.388	213.734
1174	271.286	217.921	261.817

```
1 # the method my_agent.train() from the previous cell returns a dictionary
2 # with training stats, namely:
3 # - duration of each episode during training,
4 # - return of each episode during training
5 # - the total number of training epochs at the end of each episode
6 # - the total number of steps simulated at the end of each episode
7
8 training_results.keys()
```

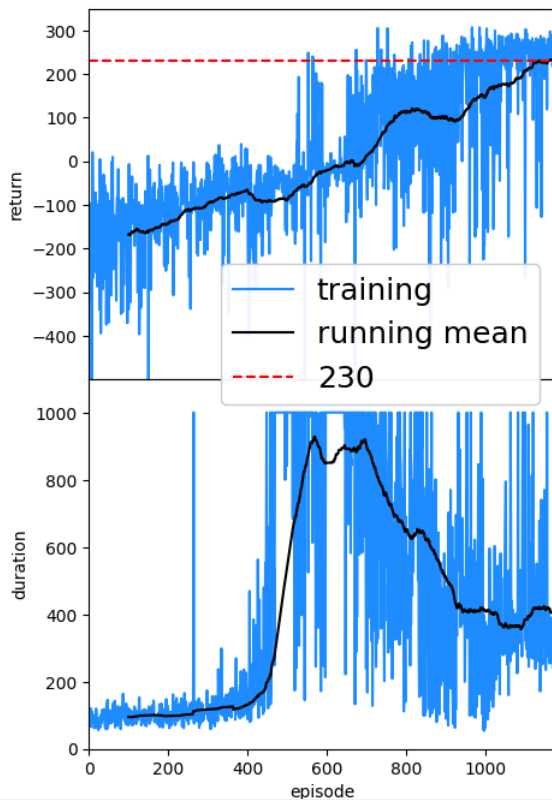
 dict_keys(['episode_durations', 'episode_returns', 'n_training_epochs', 'n_steps_simulated', 'training_completed'])

Plot training stats

```

1 # Plot both the return per episode and the duration per episode during
2 # training, together with their running average over 20 consecutive episodes
3
4 N = 100 # number of episodes for running average
5
6 def running_mean(x,N=100):
7     x_out = np.zeros(len(x)-N,dtype=float)
8     for i in range(len(x)-N):
9         x_out[i] = np.mean(x[i:i+N+1])
10    return x_out
11
12 def plot_returns_and_durations(training_results,filename=None):
13     fig.axes = plt.subplots(2,1,figsize=(5,8))
14     fig.subplots_adjust(hspace=0.0001)
15     #
16     # return as a function of episode
17     ax = axes[0]
18     x = training_results['episode_returns']
19     t = np.arange(len(x)) + 1
20     #
21     ax.plot(t,x,label='training',color='dodgerblue',)
22     # add running mean
23     x = running_mean(x=x,N=N)
24     t = np.arange(len(x)) + N
25     ax.plot(t,x,color='black',label='running mean')
26     #
27     ax.axhline(230,ls='--',
28               label='230',
29               color='red')
30     #
31     ax.set_ylim(-499,350)
32     ax.set_xticks([])
33     ax.set_xlim(0,len(t)+100)
34     ax.set_xlabel(r'episode')
35     ax.set_ylabel(r'return')
36     #
37     #
38     ax = axes[1]
39     x = training_results['episode_durations']
40     t = np.arange(len(x)) + 1
41     #
42     ax.plot(t,x,label='training',color='dodgerblue',)
43     # add running mean
44     x = running_mean(x=x,N=N)
45     t = np.arange(len(x)) + N
46     ax.plot(t,x,color='black',label='running mean')
47     #
48     ax.axhline(1200,ls='--', # draw line outside of plot scale,
49               label='230', # to get the red dotted line into the legend
50               color='red')
51     #
52     ax.set_ylim(0,1100)
53     ax.set_xlim(0,len(t)+100)
54     ax.set_xlabel(r'episode')
55     ax.set_ylabel(r'duration')
56     ax.legend(loc='upper right',bbox_to_anchor=(1.,1.35),
57             framealpha=0.95,
58             fontsize=18)
59     #
60     plt.show()
61     if filename != None:
62         fig.savefig(filename,bbox_inches='tight')
63     plt.close(fig)
64
65 plot_returns_and_durations(training_results=training_results)

```



✓ Create gameplay video using trained agent

First we create a "live" video that pops up and shows Lunar Lander gameplay performed by the agent

```
1 # There is the issue that the game window freezes when running gym games
2 # in jupyter notebooks, see https://github.com/openai/gym/issues/2433
3 # We here use the fix from that website, which is to use the following
4 # wrapper class:
5 class PyGameWrapper(gym.Wrapper):
6     def render(self, **kwargs):
7         retval = self.env.render(**kwargs)
8         for event in pygame.event.get():
9             pass
10        return retval

1 # Create a wrapped environment
2 env = PyGameWrapper(gym.make('LunarLander-v3', render_mode='human'))
3
4 N_episodes = 100
5
6 result_string = 'Run {0}: duration = {1}, total return = {2:7.3f}'
7
8 for j in range(N_episodes):
9     state, info = env.reset()
10
11    total_reward = 0
12    for i in itertools.count():
13        #env.render()
14
15        action = my_agent.act(state)
16        state, reward, terminated, truncated, info = env.step(action)
17        done = terminated or truncated
18        total_reward += reward
19
20    if done:
21        print(result_string.format(j+1, i+1, total_reward))
22        break
23
24 env.close()
```



```

Run 55: duration = 239, total return = 233.419
Run 56: duration = 494, total return = 246.716
Run 57: duration = 290, total return = 272.497
Run 58: duration = 316, total return = 268.543
Run 59: duration = 305, total return = 282.570
Run 60: duration = 266, total return = 281.784
Run 61: duration = 287, total return = 254.615
Run 62: duration = 253, total return = 258.342
Run 63: duration = 269, total return = 248.640
Run 64: duration = 258, total return = 284.505
Run 65: duration = 418, total return = 237.460
Run 66: duration = 298, total return = 256.810
Run 67: duration = 281, total return = 273.306
Run 68: duration = 241, total return = 267.006
Run 69: duration = 202, total return = 7.432
Run 70: duration = 260, total return = 301.529
Run 71: duration = 228, total return = 287.347
Run 72: duration = 241, total return = 292.207
Run 73: duration = 260, total return = 286.294
Run 74: duration = 284, total return = 297.037
Run 75: duration = 363, total return = 244.174
Run 76: duration = 227, total return = 284.547
Run 77: duration = 235, total return = 285.826
Run 78: duration = 261, total return = 271.271
Run 79: duration = 227, total return = 257.818
Run 80: duration = 274, total return = 282.016
Run 81: duration = 255, total return = 253.038
Run 82: duration = 278, total return = 260.515
Run 83: duration = 287, total return = 293.321
Run 84: duration = 351, total return = 265.774
Run 85: duration = 299, total return = 253.173
Run 86: duration = 241, total return = 243.084
Run 87: duration = 230, total return = 275.609
Run 88: duration = 270, total return = 263.973
Run 89: duration = 241, total return = 281.905
Run 90: duration = 286, total return = 256.052
Run 91: duration = 326, total return = 251.276
Run 92: duration = 380, total return = 204.303
Run 93: duration = 227, total return = 276.534
Run 94: duration = 361, total return = 282.234
Run 95: duration = 239, total return = 246.895
Run 96: duration = 287, total return = 254.795
Run 97: duration = 230, total return = 267.565
Run 98: duration = 241, total return = 258.410
Run 99: duration = 264, total return = 265.239
Run 100: duration = 271, total return = 243.953

```

We also create a video file containing 20 games played by the agent

```

1 import gymnasium as gym
2 from gymnasium.wrappers import RecordEpisodeStatistics, RecordVideo
3
4 num_eval_episodes = 100
5
6 env = gym.make("LunarLander-v3", render_mode="rgb_array") # replace with your environment
7 env = RecordVideo(env, video_folder="my_video", name_prefix="eval",
8                  episode_trigger=lambda x: True)
9
10 env = RecordEpisodeStatistics(env, buffer_length=num_eval_episodes)
11
12 for episode_num in range(num_eval_episodes):
13     obs, info = env.reset()
14
15     episode_over = False
16     while not episode_over:
17         action = my_agent.act(state)
18         #action = env.action_space.sample() # replace with actual agent
19         obs, reward, terminated, truncated, info = env.step(action)
20
21         episode_over = terminated or truncated
22 env.close()
23
24 print(f'Episode time taken: {env.time_queue}')
25 print(f'Episode total rewards: {env.return_queue}')
26 print(f'Episode lengths: {env.length_queue}')

```

```

🔍 Episode time taken: deque([0.246945, 0.239493, 0.215686, 0.243252, 0.32205, 0.353517, 0.473401, 0.398357, 0.563515, 0.219236, 0.297372, 0.210349, 0.20769, 0.2263,
Episode total rewards: deque([-129.14787369042065, -167.01613602891018, -133.04918564776568, -132.18081214539345, -291.02270670900907, -185.18266855311606, -128.5,
Episode lengths: deque([75, 73, 69, 71, 97, 76, 86, 71, 90, 69, 102, 67, 71, 75, 85, 85, 52, 84, 88, 89, 83, 56, 76, 65, 73, 66, 86, 88, 77, 62, 67, 67, 57, 60, ...])

```

Video Download

```

1 import os
2 import shutil
3 from google.colab import drive
4
5 # Google 드라이브 마운트
6 drive.mount('/content/drive')
7
8 # 원본 경로
9 source_path = '/content/20242R0136C0SE47402/FinalProject/lunar-lander/mv_video'

```

```
10
11 # 목적지 경로
12 destination_path = '/content/drive/MyDrive/lunar-lander/mine'
13
14 # 목적지 경로가 없으면 생성
15 if not os.path.exists(destination_path):
16     os.makedirs(destination_path)
17
18 # '/content' 경로에 있는 모든 파일과 폴더를 이동
19 for filename in os.listdir(source_path):
20     file_path = os.path.join(source_path, filename)
21     if os.path.isfile(file_path) or os.path.isdir(file_path):
22         shutil.move(file_path, destination_path)
```

↗ Drive already mounted at /content/drive: to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).