

▼ Homework 4

Instructions

- This homework focuses on understanding and applying CoCoOp for CLIP prompt tuning. It consists of **four questions** designed to assess both theoretical understanding and practical application.
- Please organize your answers and results for the questions below and submit this jupyter notebook as a **.pdf file**.
- **Deadline: 11/26 (Sat) 23:59**

▼ Preparation

- Run the code below before proceeding with the homework (Q1, Q2).
- If an error occurs, click 'Run Session Again' and then restart the runtime from the beginning.

```

1 !git clone https://github.com/mlvlab/ProMetaR.git
2 %cd ProMetaR/
3
4 !git clone https://github.com/KaiyangZhou/Dassl.pytorch.git
5 %cd Dassl.pytorch/
6
7 # Install dependencies
8 !pip install -r requirements.txt
9 !cp -r dassl ../
10 # Install this library (no need to re-build if the source code is modified)
11 # !python setup.py develop
12 %cd ..
13
14 !pip install -r requirements.txt
15
16 %mkdir outputs
17 %mkdir data
18
19 %cd data
20 %mkdir eurosat
21 !wget http://madm.dfki.de/files/sentinel/EuroSAT.zip EuroSAT.zip
22
23 !unzip -o EuroSAT.zip -d eurosat/
24 %cd eurosat
25 !gdown 1lp7yaCWF10ea0FUGga0lUdVi_DDQth1o
26
27 %cd ../../
28
29 import os.path as osp
30 from collections import OrderedDict
31 import math
32 import torch
33 import torch.nn as nn
34 from torch.nn import functional as F
35 from torch.cuda.amp import GradScaler, autocast
36 from PIL import Image
37 import torchvision.transforms as transforms
38 import torch
39 from clip import clip
40 from clip.simple_tokenizer import SimpleTokenizer as _Tokenizer
41 import time
42 from tqdm import tqdm
43 import datetime
44 import argparse
45 from dassl.utils import setup_logger, set_random_seed, collect_env_info
46 from dassl.config import get_cfg_default
47 from dassl.engine import build_trainer
48 from dassl.engine import TRAINER_REGISTRY, TrainerX
49 from dassl.metrics import compute_accuracy
50 from dassl.utils import load_pretrained_weights, load_checkpoint
51 from dassl.optim import build_optimizer, build_lr_scheduler
52
53 # custom
54 import datasets.oxford_pets
55 import datasets.oxford_flowers
56 import datasets.fgvc_aircraft
57 import datasets.dtd
58 import datasets.eurosat
59 import datasets.stanford_cars
60 import datasets.food101
61 import datasets.sun397
62 import datasets.caltech101
63 import datasets.ucf101
64 import datasets.imagenet
65 import datasets.imagenet_sketch
66 import datasets.imagenetv2
67 import datasets.imagenet_a
68 import datasets.imagenet_r
69
70 def print_args(args, cfg):

```

```

71     print("*****")
72     print("** Arguments **")
73     print("*****")
74     optkeys = list(args.__dict__.keys())
75     optkeys.sort()
76     for key in optkeys:
77         print("{}: {}".format(key, args.__dict__[key]))
78     print("*****")
79     print("** Config **")
80     print("*****")
81     print(cfg)
82
83 def reset_cfg(cfg, args):
84     if args.root:
85         cfg.DATASET.ROOT = args.root
86     if args.output_dir:
87         cfg.OUTPUT_DIR = args.output_dir
88     if args.seed:
89         cfg.SEED = args.seed
90     if args.trainer:
91         cfg.TRAINER.NAME = args.trainer
92     cfg.DATASET.NUM_SHOTS = 16
93     cfg.DATASET.SUBSAMPLE_CLASSES = args.subsample_classes
94     cfg.DATALOADER.TRAIN_X.BATCH_SIZE = args.train_batch_size
95     cfg.OPTIM.MAX_EPOCH = args.epoch
96
97 def extend_cfg(cfg):
98     """
99     Add new config variables.
100    """
101    from yacs.config import CfgNode as CN
102    cfg.TRAINER.COOP = CN()
103    cfg.TRAINER.COOP.N_CTX = 16 # number of context vectors
104    cfg.TRAINER.COOP.CSC = False # class-specific context
105    cfg.TRAINER.COOP.CTX_INIT = "" # initialization words
106    cfg.TRAINER.COOP.PREC = "fp16" # fp16, fp32, amp
107    cfg.TRAINER.COOP.CLASS_TOKEN_POSITION = "end" # 'middle' or 'end' or 'front'
108    cfg.TRAINER.COOCOOP = CN()
109    cfg.TRAINER.COOCOOP.N_CTX = 4 # number of context vectors
110    cfg.TRAINER.COOCOOP.CTX_INIT = "a photo of a" # initialization words
111    cfg.TRAINER.COOCOOP.PREC = "fp16" # fp16, fp32, amp
112    cfg.TRAINER.PROMETAR = CN()
113    cfg.TRAINER.PROMETAR.N_CTX_VISION = 4 # number of context vectors at the vision branch
114    cfg.TRAINER.PROMETAR.N_CTX_TEXT = 4 # number of context vectors at the language branch
115    cfg.TRAINER.PROMETAR.CTX_INIT = "a photo of a" # initialization words
116    cfg.TRAINER.PROMETAR.PREC = "fp16" # fp16, fp32, amp
117    cfg.TRAINER.PROMETAR.PROMPT_DEPTH_VISION = 9 # Max 12, minimum 0, for 0 it will be using shallow IVLP prompting (J=1)
118    cfg.TRAINER.PROMETAR.PROMPT_DEPTH_TEXT = 9 # Max 12, minimum 0, for 0 it will be using shallow IVLP prompting (J=1)
119    cfg.DATASET.SUBSAMPLE_CLASSES = "all" # all, base or new
120    cfg.TRAINER.PROMETAR.ADAPT_LR = 0.0005
121    cfg.TRAINER.PROMETAR.LR_RATIO = 0.0005
122    cfg.TRAINER.PROMETAR.FAST_ADAPTATION = False
123    cfg.TRAINER.PROMETAR.MIXUP_ALPHA = 0.5
124    cfg.TRAINER.PROMETAR.MIXUP_BETA = 0.5
125    cfg.TRAINER.PROMETAR.DIM_RATE=8
126    cfg.OPTIM_VNET = CN()
127    cfg.OPTIM_VNET.NAME = "adam"
128    cfg.OPTIM_VNET.LR = 0.0003
129    cfg.OPTIM_VNET.WEIGHT_DECAY = 5e-4
130    cfg.OPTIM_VNET.MOMENTUM = 0.9
131    cfg.OPTIM_VNET.SGD_DAMPNING = 0
132    cfg.OPTIM_VNET.SGD_NESTEROV = False
133    cfg.OPTIM_VNET.RMSPROP_ALPHA = 0.99
134    cfg.OPTIM_VNET.ADM_BETA1 = 0.9
135    cfg.OPTIM_VNET.ADM_BETA2 = 0.999
136    cfg.OPTIM_VNET.STAGED_LR = False
137    cfg.OPTIM_VNET.NEW_LAYERS = ()
138    cfg.OPTIM_VNET.BASE_LR_MULT = 0.1
139    # Learning rate scheduler
140    cfg.OPTIM_VNET.LR_SCHEDULER = "single_step"
141    # -1 or 0 means the stepsize is equal to max_epoch
142    cfg.OPTIM_VNET.STEPSIZE = (-1, )
143    cfg.OPTIM_VNET.GAMMA = 0.1
144    cfg.OPTIM_VNET.MAX_EPOCH = 10
145    # Set WARMUP_EPOCH larger than 0 to activate warmup training
146    cfg.OPTIM_VNET.WARMUP_EPOCH = -1
147    # Either linear or constant
148    cfg.OPTIM_VNET.WARMUP_TYPE = "linear"
149    # Constant learning rate when type=constant
150    cfg.OPTIM_VNET.WARMUP_CONS_LR = 1e-5
151    # Minimum learning rate when type=linear
152    cfg.OPTIM_VNET.WARMUP_MIN_LR = 1e-5
153    # Recount epoch for the next scheduler (last_epoch=-1)
154    # Otherwise last_epoch=warmup_epoch
155    cfg.OPTIM_VNET.WARMUP_RECOUNT = True
156
157 def setup_cfg(args):
158     cfg = get_cfg_default()
159     extend_cfg(cfg)
160     # 1. From the dataset config file
161     if args.dataset_config_file:

```

```

162         cfg.merge_from_file(args.dataset_config_file)
163     # 2. From the method config file
164     if args.config_file:
165         cfg.merge_from_file(args.config_file)
166     # 3. From input arguments
167     reset_cfg(cfg, args)
168     cfg.freeze()
169     return cfg
170
171 _tokenizer = _Tokenizer()
172
173 def load_clip_to_cpu(cfg): # Load CLIP
174     backbone_name = cfg.MODEL.BACKBONE.NAME
175     url = clip._MODELS[backbone_name]
176     model_path = clip._download(url)
177
178     try:
179         # loading JIT archive
180         model = torch.jit.load(model_path, map_location="cpu").eval()
181         state_dict = None
182
183     except RuntimeError:
184         state_dict = torch.load(model_path, map_location="cpu")
185
186     if cfg.TRAINER.NAME == "":
187         design_trainer = "CoOp"
188     else:
189         design_trainer = cfg.TRAINER.NAME
190     design_details = {"trainer": design_trainer,
191                      "vision_depth": 0,
192                      "language_depth": 0, "vision_ctx": 0,
193                      "language_ctx": 0}
194     model = clip.build_model(state_dict or model.state_dict(), design_details)
195
196     return model
197
198 from dassl.config import get_cfg_default
199 cfg = get_cfg_default()
200 cfg.MODEL.BACKBONE.NAME = "ViT-B/16" # Set the vision encoder backbone of CLIP to ViT.
201 clip_model = load_clip_to_cpu(cfg)
202
203
204
205 class TextEncoder(nn.Module):
206     def __init__(self, clip_model): # 초기화 하는 함수
207         super().__init__()
208         self.transformer = clip_model.transformer
209         self.positional_embedding = clip_model.positional_embedding
210         self.ln_final = clip_model.ln_final
211         self.text_projection = clip_model.text_projection
212         self.dtype = clip_model.dtype
213
214     def forward(self, prompts, tokenized_prompts): # 모델 호출
215         x = prompts + self.positional_embedding.type(self.dtype)
216         x = x.permute(1, 0, 2) # NLD -> LND
217         x = self.transformer(x)
218         x = x.permute(1, 0, 2) # LND -> NLD
219         x = self.ln_final(x).type(self.dtype)
220
221         # x.shape = [batch_size, n_ctx, transformer.width]
222         # take features from the eot embedding (eot_token is the highest number in each sequence)
223         x = x[torch.arange(x.shape[0]), tokenized_prompts.argmax(dim=-1)] @ self.text_projection
224
225         return x
226
227
228 @TRAINER_REGISTRY.register(force=True)
229 class CoCoOp(TrainerX):
230     def check_cfg(self, cfg):
231         assert cfg.TRAINER.COOCOOP.PREC in ["fp16", "fp32", "amp"]
232
233     def build_model(self):
234         cfg = self.cfg
235         classnames = self.dm.dataset.classnames
236         print(f"Loading CLIP (backbone: {cfg.MODEL.BACKBONE.NAME})")
237         clip_model = load_clip_to_cpu(cfg)
238
239         if cfg.TRAINER.COOCOOP.PREC == "fp32" or cfg.TRAINER.COOCOOP.PREC == "amp":
240             # CLIP's default precision is fp16
241             clip_model.float()
242
243         print("Building custom CLIP")
244         self.model = CoCoOpCustomCLIP(cfg, classnames, clip_model)
245
246         print("Turning off gradients in both the image and the text encoder")
247         name_to_update = "prompt_learner"
248
249         for name, param in self.model.named_parameters():
250             if name_to_update not in name:
251                 param.requires_grad_(False)
252

```

```

253     # Double check
254     enabled = set()
255     for name, param in self.model.named_parameters():
256         if param.requires_grad:
257             enabled.add(name)
258     print(f"Parameters to be updated: {enabled}")
259
260     if cfg.MODEL.INIT_WEIGHTS:
261         load_pretrained_weights(self.model.prompt_learner, cfg.MODEL.INIT_WEIGHTS)
262
263     self.model.to(self.device)
264     # NOTE: only give prompt_learner to the optimizer
265     self.optim = build_optimizer(self.model.prompt_learner, cfg.OPTIM)
266     self.sched = build_lr_scheduler(self.optim, cfg.OPTIM)
267     self.register_model("prompt_learner", self.model.prompt_learner, self.optim, self.sched)
268
269     self.scaler = GradScaler() if cfg.TRAINER.COCOOOP.PREC == "amp" else None
270
271     # Note that multi-gpu training could be slow because CLIP's size is
272     # big, which slows down the copy operation in DataParallel
273     device_count = torch.cuda.device_count()
274     if device_count > 1:
275         print(f"Multiple GPUs detected (n_gpus={device_count}), use all of them!")
276         self.model = nn.DataParallel(self.model)
277
278     def before_train(self):
279         directory = self.cfg.OUTPUT_DIR
280         if self.cfg.RESUME:
281             directory = self.cfg.RESUME
282         self.start_epoch = self.resume_model_if_exist(directory)
283
284         # Remember the starting time (for computing the elapsed time)
285         self.time_start = time.time()
286
287     def forward_backward(self, batch):
288         image, label = self.parse_batch_train(batch)
289
290         model = self.model
291         optim = self.optim
292         scaler = self.scaler
293
294         prec = self.cfg.TRAINER.COCOOOP.PREC
295         loss = model(image, label) # Input image 모델 통과
296         optim.zero_grad()
297         loss.backward() # Backward (역전파)
298         optim.step() # 모델 parameter update
299
300         loss_summary = {"loss": loss.item()}
301
302         if (self.batch_idx + 1) == self.num_batches:
303             self.update_lr()
304
305         return loss_summary
306
307     def parse_batch_train(self, batch):
308         input = batch["img"]
309         label = batch["label"]
310         input = input.to(self.device)
311         label = label.to(self.device)
312         return input, label
313
314     def load_model(self, directory, epoch=None):
315         if not directory:
316             print("Note that load_model() is skipped as no pretrained model is given")
317             return
318
319         names = self.get_model_names()
320
321         # By default, the best model is loaded
322         model_file = "model-best.pth.tar"
323
324         if epoch is not None:
325             model_file = "model.pth.tar-" + str(epoch)
326
327         for name in names:
328             model_path = osp.join(directory, name, model_file)
329
330             if not osp.exists(model_path):
331                 raise FileNotFoundError('Model not found at "{}".format(model_path)')
332
333             checkpoint = load_checkpoint(model_path)
334             state_dict = checkpoint["state_dict"]
335             epoch = checkpoint["epoch"]
336
337             # Ignore fixed token vectors
338             if "token_prefix" in state_dict:
339                 del state_dict["token_prefix"]
340
341             if "token_suffix" in state_dict:
342                 del state_dict["token_suffix"]
343

```

```

344
345         print("Loading weights to {} " 'from "{}" (epoch = {})'.format(name, model_path, epoch))
346         # set strict=False
347         self._models[name].load_state_dict(state_dict, strict=False)
348
349     def after_train(self):
350         print("Finish training")
351
352     do_test = not self.cfg.TEST.NO_TEST
353     if do_test:
354         if self.cfg.TEST.FINAL_MODEL == "best_val":
355             print("Deploy the model with the best val performance")
356             self.load_model(self.output_dir)
357         else:
358             print("Deploy the last-epoch model")
359         acc = self.test()
360
361     # Show elapsed time
362     elapsed = round(time.time() - self.time_start)
363     elapsed = str(datetime.timedelta(seconds=elapsed))
364     print(f"Elapsed: {elapsed}")
365
366     # Close writer
367     self.close_writer()
368     return acc
369
370     def train(self):
371         """Generic training loops."""
372         self.before_train()
373         for self.epoch in range(self.start_epoch, self.max_epoch):
374             self.before_epoch()
375             self.run_epoch()
376             self.after_epoch()
377         acc = self.after_train()
378         return acc
379
380     parser = argparse.ArgumentParser()
381     parser.add_argument("--root", type=str, default="data/", help="path to dataset")
382     parser.add_argument("--output-dir", type=str, default="outputs/cocoop3", help="output directory")
383     parser.add_argument(
384         "--seed", type=int, default=1, help="only positive value enables a fixed seed"
385     )
386     parser.add_argument(
387         "--config-file", type=str, default="configs/trainers/ProMetaR/vit_b16_c2_ep10_batch4_4+4ctx.yaml", help="path to config file"
388     )
389     parser.add_argument(
390         "--dataset-config-file",
391         type=str,
392         default="configs/datasets/eurosat.yaml",
393         help="path to config file for dataset setup",
394     )
395     parser.add_argument("--trainer", type=str, default="CoOp", help="name of trainer")
396     parser.add_argument("--eval-only", action="store_true", help="evaluation only")
397     parser.add_argument(
398         "--model-dir",
399         type=str,
400         default="",
401         help="load model from this directory for eval-only mode",
402     )
403     parser.add_argument("--train-batch-size", type=int, default=4)
404     parser.add_argument("--epoch", type=int, default=10)
405     parser.add_argument("--subsample-classes", type=str, default="base")
406     parser.add_argument(
407         "--load-epoch", type=int, default=0, help="load model weights at this epoch for evaluation"
408     )
409     args = parser.parse_args([])
410
411     def main(args):
412         cfg = setup_cfg(args)
413         if cfg.SEED >= 0:
414             set_random_seed(cfg.SEED)
415
416         if torch.cuda.is_available() and cfg.USE_CUDA:
417             torch.backends.cudnn.benchmark = True
418
419         trainer = build_trainer(cfg)
420         if args.eval_only:
421             trainer.load_model(args.model_dir, epoch=args.load_epoch)
422             acc = trainer.test()
423             return acc
424
425         acc = trainer.train()
426         return acc

```



```

inflating: eurosat/2750/PermanentCrop/PermanentCrop_2015.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_828.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_1106.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_1670.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_1211.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_2304.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_273.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_1088.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_612.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_1438.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_164.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_1059.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_505.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_977.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_2475.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_1912.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_1560.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_2014.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_1101.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_1677.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_19.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_1216.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_2303.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_1753.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_1332.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_1495.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_2227.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_118.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_1444.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_1836.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_2130.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_1782.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_579.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_1025.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_2409.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_853.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_421.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_386.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_2068.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_882.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_357.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_1.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_65.jpg
inflating: eurosat/2750/PermanentCrop/PermanentCrop_736.jpg
/content/ProMetaR/data/eurosat
Downloading...
From: https://drive.google.com/uc?id=1lp7yaCWF10ea0FUGga01UdVi\_DDQth1o
To: /content/ProMetaR/data/eurosat/split_zhou-EuroSAT.json
100% 3.01M/3.01M [00:00<00:00, 128MB/s]
/content/ProMetaR

```

▼ Q1. Understanding and implementing CoCoOp

- We have learned how to define CoOp in Lab Session 4.
- The main difference between CoOp and CoCoOp is **meta network** to extract image tokens that is added to the text prompt.
- Based on the CoOp code given in Lab Session 4, fill-in-the-blank exercise to test your understanding of critical parts of the CoCoOp.

```

1 import torch.nn as nn
2
3 class CoCoOpPromptLearner(nn.Module):
4     def __init__(self, cfg, classnames, clip_model):
5         super().__init__()
6         n_cls = len(classnames)
7         n_ctx = cfg.TRAINER.COOCOOP.N_CTX
8         ctx_init = cfg.TRAINER.COOCOOP.CTX_INIT
9         dtype = clip_model.dtype
10        ctx_dim = clip_model.ln_final.weight.shape[0]
11        vis_dim = clip_model.visual.output_dim
12        clip_imsize = clip_model.visual.input_resolution
13        cfg_imsize = cfg.INPUT.SIZE[0]
14        assert cfg_imsize == clip_imsize, f"cfg_imsize ({cfg_imsize}) must equal to clip_imsize ({clip_imsize})"
15
16        if ctx_init:
17            # use given words to initialize context vectors
18            ctx_init = ctx_init.replace("_", " ")
19            n_ctx = len(ctx_init.split(" "))
20            prompt = clip.tokenize(ctx_init)
21            with torch.no_grad():
22                embedding = clip_model.token_embedding(prompt).type(dtype)
23            ctx_vectors = embedding[0, 1: 1 + n_ctx, :]
24            prompt_prefix = ctx_init
25        else:
26            # random initialization
27            ctx_vectors = torch.empty(n_ctx, ctx_dim, dtype=dtype)
28            nn.init.normal_(ctx_vectors, std=0.02)
29            prompt_prefix = " ".join(["X"] * n_ctx)
30
31        print(f'Initial context: "{prompt_prefix}")
32        print(f'Number of context words (tokens): {n_ctx}')
33
34        self.ctx = nn.Parameter(ctx_vectors) # Wrap the initialized prompts above as parameters to make them trainable.
35
36        ### Tokenize ###

```

```

37     classnames = [name.replace("_", " ") for name in classnames] # 예) "Forest"
38     name_lens = [len(_tokenizer.encode(name)) for name in classnames]
39     prompts = [prompt_prefix + " " + name + "." for name in classnames] # 예) "A photo of Forest."
40
41     tokenized_prompts = torch.cat([clip.tokenize(p) for p in prompts]) # 예) [49406, 320, 1125, 539...]
42
43
44
45     #####
46     ##### Q1. Fill in the blank #####
47     ##### Define Meta Net #####
48     self.meta_net = nn.Sequential(OrderedDict([
49         ("linear1", nn.Linear(vis_dim, vis_dim // 16)),
50         ("relu", nn.ReLU(inplace=True)),
51         ("linear2", nn.Linear(vis_dim // 16, ctx_dim))
52     ]))
53     #####
54     ## Hint: meta network is composed to linear layer, relu activation, and linear layer.
55
56
57
58     if cfg.TRAINER.COOCOOP.PREC == "fp16":
59         self.meta_net.half()
60
61     with torch.no_grad():
62         embedding = clip_model.token_embedding(tokenized_prompts).type(dtype)
63
64     # These token vectors will be saved when in save_model(),
65     # but they should be ignored in load_model() as we want to use
66     # those computed using the current class names
67     self.register_buffer("token_prefix", embedding[:, :1, :]) # SOS
68     self.register_buffer("token_suffix", embedding[:, 1 + n_ctx:, :]) # CLS, EOS
69     self.n_cls = n_cls
70     self.n_ctx = n_ctx
71     self.tokenized_prompts = tokenized_prompts # torch.Tensor
72     self.name_lens = name_lens
73
74     def construct_prompts(self, ctx, prefix, suffix, label=None):
75         # dim0 is either batch_size (during training) or n_cls (during testing)
76         # ctx: context tokens, with shape of (dim0, n_ctx, ctx_dim)
77         # prefix: the sos token, with shape of (n_cls, 1, ctx_dim)
78         # suffix: remaining tokens, with shape of (n_cls, *, ctx_dim)
79
80         if label is not None:
81             prefix = prefix[label]
82             suffix = suffix[label]
83
84         prompts = torch.cat(
85             [
86                 prefix, # (dim0, 1, dim)
87                 ctx, # (dim0, n_ctx, dim)
88                 suffix, # (dim0, *, dim)
89             ],
90             dim=1,
91         )
92
93         return prompts
94
95     def forward(self, im_features):
96         prefix = self.token_prefix
97         suffix = self.token_suffix
98         ctx = self.ctx # (n_ctx, ctx_dim)
99
100
101
102     #####
103     ##### Q2,3. Fill in the blank #####
104     # Hint: Image feature is given as input to meta network
105     bias = self.meta_net(im_features) # (batch, ctx_dim)
106     bias = bias.unsqueeze(1) # (batch, 1, ctx_dim)
107     ctx = ctx.unsqueeze(0) # (1, n_ctx, ctx_dim)
108     # Hint: Add meta token to context token
109     ctx_shifted = ctx + bias # (batch, n_ctx, ctx_dim)
110     #####
111     #####
112
113
114
115     # Use instance-conditioned context tokens for all classes
116     prompts = []
117     for ctx_shifted_i in ctx_shifted:
118         ctx_i = ctx_shifted_i.unsqueeze(0).expand(self.n_cls, -1, -1)
119         pts_i = self.construct_prompts(ctx_i, prefix, suffix) # (n_cls, n_tkn, ctx_dim)
120         prompts.append(pts_i)
121     prompts = torch.stack(prompts)
122
123     return prompts

```

```

1 class CoCoOpCustomCLIP(nn.Module):
2     def __init__(self, cfg, classnames, clip_model):
3         super().__init__()

```

```

4     self.prompt_learner = CoCoOpPromptLearner(cfg, classnames, clip_model)
5     self.tokenized_prompts = self.prompt_learner.tokenized_prompts
6     self.image_encoder = clip_model.visual
7     self.text_encoder = TextEncoder(clip_model)
8     self.logit_scale = clip_model.logit_scale
9     self.dtype = clip_model.dtype
10
11     def forward(self, image, label=None):
12         tokenized_prompts = self.tokenized_prompts
13         logit_scale = self.logit_scale.exp()
14
15         image_features = self.image_encoder(image.type(self.dtype))
16         image_features = image_features / image_features.norm(dim=-1, keepdim=True)
17
18         #####
19         ##### Q4. Fill in the blank #####
20         prompts = self.prompt_learner(image_features)
21         #####
22         #####
23         #####
24
25
26         logits = []
27         for pts_i, imf_i in zip(prompts, image_features):
28             text_features = self.text_encoder(pts_i, tokenized_prompts)
29             text_features = text_features / text_features.norm(dim=-1, keepdim=True)
30             l_i = logit_scale * imf_i @ text_features.t()
31             logits.append(l_i)
32         logits = torch.stack(logits)
33
34         if self.prompt_learner.training:
35             return F.cross_entropy(logits, label)
36
37         return logits

```

Q2. Training CoCoOp

In this task, you will train CoCoOp on the EuroSAT dataset. If your implementation of CoCoOp in Question 1 is correct, the following code should execute without errors. Please submit the execution file so we can evaluate whether your code runs without any issues.

```

1 # Train on the Base Classes Train split and evaluate accuracy on the Base Classes Test split.
2 args.trainer = "CoCoOp"
3 args.train_batch_size = 4
4 args.epoch = 100
5 args.output_dir = "outputs/cocoop"
6
7 args.subsample_classes = "base"
8 args.eval_only = False
9 cocoop_base_acc = main(args)

```




```

epoch [99/100] batch [20/20] time 0.100 (0.100) data 0.000 (0.030) loss 0.0691 (0.1264) lr 2.4666e-06 eta 0:00:03
epoch [100/100] batch [20/20] time 0.101 (0.127) data 0.000 (0.020) loss 0.0025 (0.1101) lr 6.1680e-07 eta 0:00:00
Checkpoint saved to outputs/cocoop/prompt_learner/model.pth.tar-100
Finish training
Deploy the last-epoch model
Evaluate on the *test* set
100%|██████████| 42/42 [01:03<00:00, 1.51s/it]> result
* total: 4,200
* correct: 3,813
* accuracy: 90.8%
* error: 9.2%
* macro_f1: 90.9%
Elapsed: 0:06:20

1 # Accuracy on the New Classes.
2 args.model_dir = "outputs/cocoop"
3 args.output_dir = "outputs/cocoop/new_classes"
4 args.subsample_classes = "new"
5 args.load_epoch = 100
6 args.eval_only = True
7 coop_novel_acc = main(args)

Loading trainer: CoCoOp
Loading dataset: EuroSAT
Reading split from /content/ProMetaR/data/eurosat/split_zhou_EuroSAT.json
Loading preprocessed few-shot data from /content/ProMetaR/data/eurosat/split_fewshot/shot_16-seed_1.pkl
SUBSAMPLE NEW CLASSES!
Building transform_train
+ random resized crop (size=(224, 224), scale=(0.08, 1.0))
+ random flip
+ to torch tensor of range [0, 1]
+ normalization (mean=[0.48145466, 0.4578275, 0.40821073], std=[0.26862954, 0.26130258, 0.27577711])
Building transform_test
+ resize the smaller edge to 224
+ 224x224 center crop
+ to torch tensor of range [0, 1]
+ normalization (mean=[0.48145466, 0.4578275, 0.40821073], std=[0.26862954, 0.26130258, 0.27577711])

-----
Dataset      EuroSAT
# classes    5
# train_x    80
# val        20
# test       3,900
-----

Loading CLIP (backbone: ViT-B/16)
/usr/local/lib/python3.10/dist-packages/torch/utils/data/dataloader.py:617: UserWarning: This DataLoader will create 8 worker processes in total. Our suggested m
warnings.warn(
/usr/local/lib/python3.10/dist-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to access the
warnings.warn(
/content/ProMetaR/dassl/utils/torchtools.py:102: FutureWarning: You are using `torch.load` with `weights_only=False` (the current default value), which uses the c
checkpoint = torch.load(fpath, map_location=map_location)
Building custom CLIP
Initial context: "a photo of a"
Number of context words (tokens): 4
Turning off gradients in both the image and the text encoder
Parameters to be updated: {'prompt_learner.meta_net.linear1.weight', 'prompt_learner.meta_net.linear1.bias', 'prompt_learner.meta_net.linear2.weight', 'prompt_lear
Loading evaluator: Classification
Loading weights to prompt_learner from "outputs/cocoop/prompt_learner/model.pth.tar-100" (epoch = 100)
Evaluate on the *test* set
100%|██████████| 39/39 [01:02<00:00, 1.60s/it]> result
* total: 3,900
* correct: 1,687
* accuracy: 43.3%
* error: 56.7%
* macro_f1: 39.0%

```

Q3. Analyzing the results of CoCoOp

Compare the results of CoCoOp with those of CoOp that we trained in Lab Session 4. Discuss possible reasons for the performance differences observed between CoCoOp and CoOp.

결과 비교

- CoOp의 train accuracy : 91.4% , test accuracy : 51.46%
- CoCoOp의 train accuracy : 90.8% , test accuracy : 43.3%

CoCoOp는 CoOp에 비해 train, test accuracy가 모두 더 낮다. 이는 기존 CoOp는 학습 가능한 프롬프트를 통해 작동하는 것과 달리, CoCoOp는 더 유연하고 민감한 prompt에 초점을 맞추었기에 발생할 수 있다. 즉, CoCoOp가 학습 데이터와의 fitting보다는 **다양한 동적 prompt**를 목표로 설계되었기 때문일 가능성이 있다.

특히 CoCoOp는 CoOp와 다르게 meta-network를 이용하여 동적 prompt를 만들고 이를 통해 도메인 간 차이에 더욱 효과적인 기능을 수행하려고 하였다. 하지만 오히려 이 과정이 test data에서의 성능 저하 및 비효율성을 야기할 수 있었다.

